

DEEP NEURAL NETWORK

Import TensorFlow

```
1 import tensorflow as tf
2
3 from tensorflow.keras import datasets, layers, models
4 import matplotlib.pyplot as plt
```

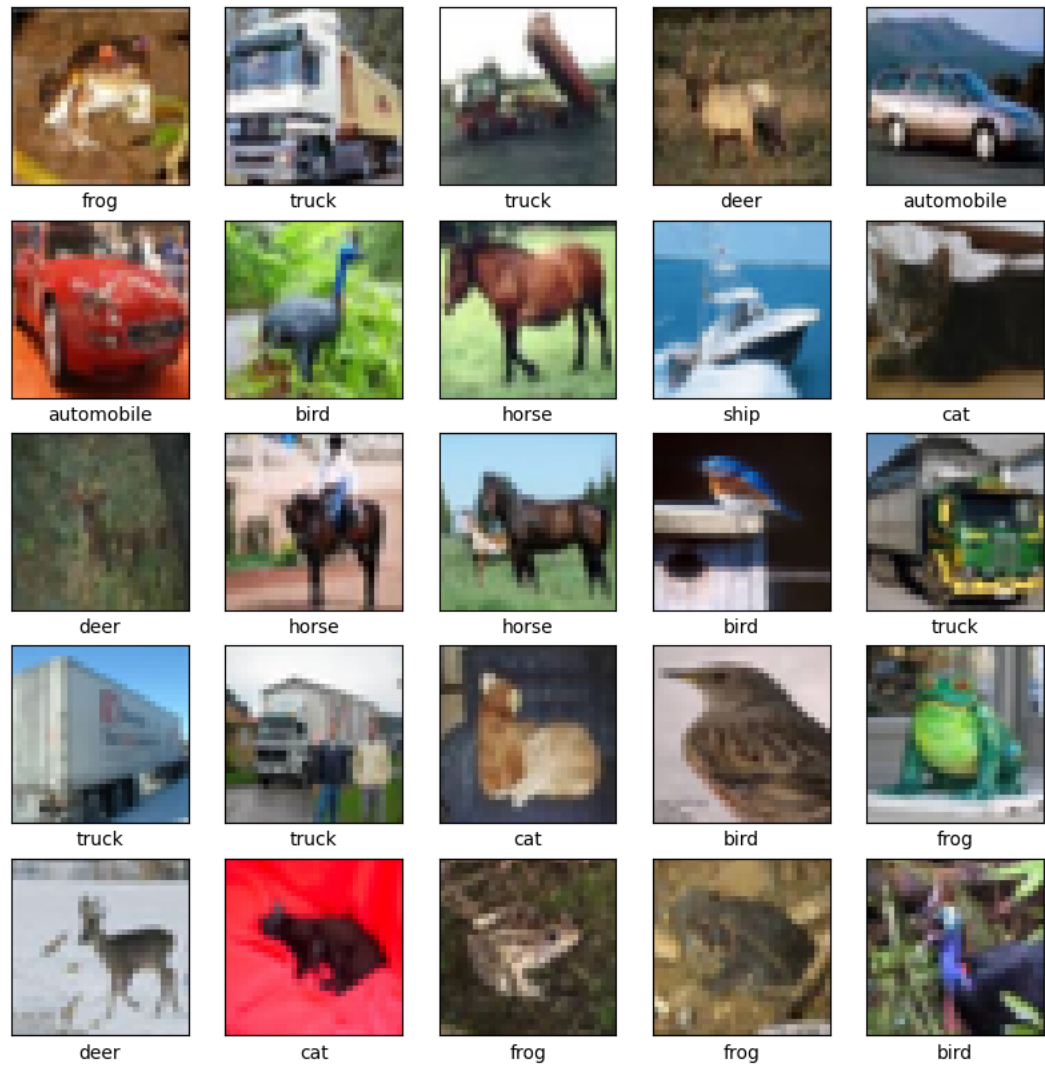
Download and prepare the CIFAR10 dataset

```
1 (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
2
3 # Normalize pixel values to be between 0 and 1
4 train_images, test_images = train_images / 255.0, test_images / 255.0

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 14s 0us/step
```

Verify the data

```
1 class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
2               'dog', 'frog', 'horse', 'ship', 'truck']
3
4 plt.figure(figsize=(10,10))
5 for i in range(25):
6     plt.subplot(5,5,i+1)
7     plt.xticks([])
8     plt.yticks([])
9     plt.grid(False)
10    plt.imshow(train_images[i])
11    # The CIFAR labels happen to be arrays,
12    # which is why you need the extra index
13    plt.xlabel(class_names[train_labels[i][0]])
14 plt.show()
```



Create the convolutional base

```
1 model = models.Sequential()
2 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
3 model.add(layers.MaxPooling2D((2, 2)))
4 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
5 model.add(layers.MaxPooling2D((2, 2)))
6 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Let's display the architecture of your model so far:

```
1 model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
-----
conv2d (Conv2D)              (None, 30, 30, 32)         896

max_pooling2d (MaxPooling2D) (None, 15, 15, 32)         0

conv2d_1 (Conv2D)            (None, 13, 13, 64)         18496

max_pooling2d_1 (MaxPooling2D) (None, 6, 6, 64)          0

conv2d_2 (Conv2D)            (None, 4, 4, 64)          36928

Total params: 56320 (220.00 KB)
Trainable params: 56320 (220.00 KB)
Non-trainable params: 0 (0.00 Byte)
```

Add Dense layers on top

```
1 model.add(layers.Flatten())
2 model.add(layers.Dense(64, activation='relu'))
3 model.add(layers.Dense(10))
```

Here's the complete architecture of your model:

```
1 model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
-----
conv2d (Conv2D)              (None, 30, 30, 32)         896

max_pooling2d (MaxPooling2D) (None, 15, 15, 32)         0

conv2d_1 (Conv2D)            (None, 13, 13, 64)         18496

max_pooling2d_1 (MaxPooling2D) (None, 6, 6, 64)          0

conv2d_2 (Conv2D)            (None, 4, 4, 64)          36928

flatten (Flatten)            (None, 1024)                0

dense (Dense)                (None, 64)                  65600

dense_1 (Dense)              (None, 10)                  650

Total params: 122570 (478.79 KB)
Trainable params: 122570 (478.79 KB)
Non-trainable params: 0 (0.00 Byte)
```

Compile and train the model

```
1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4
5 history = model.fit(train_images, train_labels, epochs=10,
6                     validation_data=(test_images, test_labels))


Epoch 1/10
1563/1563 [=====] - 20s 6ms/step - loss: 1.5405 - accuracy: 0.4390 - val_loss: 1.2710 - val_accuracy: 0.5444
Epoch 2/10
1563/1563 [=====] - 8s 5ms/step - loss: 1.1881 - accuracy: 0.5788 - val_loss: 1.0931 - val_accuracy: 0.6108
Epoch 3/10
1563/1563 [=====] - 9s 6ms/step - loss: 1.0436 - accuracy: 0.6329 - val_loss: 0.9882 - val_accuracy: 0.6583
Epoch 4/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.9511 - accuracy: 0.6669 - val_loss: 1.0265 - val_accuracy: 0.6377
Epoch 5/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.8768 - accuracy: 0.6919 - val_loss: 0.9480 - val_accuracy: 0.6696
Epoch 6/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.8200 - accuracy: 0.7114 - val_loss: 0.9161 - val_accuracy: 0.6854
Epoch 7/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.7691 - accuracy: 0.7302 - val_loss: 0.8694 - val_accuracy: 0.7044
Epoch 8/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.7235 - accuracy: 0.7471 - val_loss: 0.8892 - val_accuracy: 0.6919
Epoch 9/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.6879 - accuracy: 0.7604 - val_loss: 0.8709 - val_accuracy: 0.7072
Epoch 10/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.6495 - accuracy: 0.7727 - val_loss: 0.8733 - val_accuracy: 0.7075
```

Evaluate the model

```
1 plt.plot(history.history['accuracy'], label='accuracy')
2 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
3 plt.xlabel('Epoch')
4 plt.ylabel('Accuracy')
5 plt.ylim([0.5, 1])
```

```
6 plt.legend(loc='lower right')
7
8 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)

313/313 - 1s - loss: 0.8733 - accuracy: 0.7075 - 691ms/epoch - 2ms/step
```




Print accuracy

```
1 print(test_acc)

0.7074999809265137
```

Prediction

```
1 n=111 ##Number of image
2
3 plt.figure(figsize=(2,2))
4 plt.imshow(test_images[n])
5 plt.xlabel(class_names[test_labels[n][0]])
6 plt.show()
```



```
1 import numpy as np
2
3
4 predictions = model.predict(test_images)
5 print(predictions[n])
6
7 print(
8     "This image most likely belongs to {} with a {:.2f} percent confidence."
9     .format(class_names[np.argmax(predictions[n])], 10 * np.max(predictions[n])))
10 )

313/313 [=====] - 1s 2ms/step
[ 6.858806  4.0483427  2.6050212 -0.4147426 -3.392217 -7.305912
 -5.7308545 -5.454667  1.9584597 -1.1152903]
This image most likely belongs to airplane with a 68.59 percent confidence.
```